

WRITE SOMETHING CLEVER

951926

1001231

1024072

1028907

Abstract—In this report, we present a system which performs the 1996 AAAI Mobile Robotics Competition task “Call a meeting”[1]. The robot was tasked with bringing a certain number of people to a meeting room previously determined as available. We use probabilistic road mapping, Monte Carlo localisation and a tripod-mounted Kinect with a face detection module running on a Pioneer 2DX platform to complete the task. In the first section, we provide some background information about the areas of robotics that are relevant to the problem. We then provide a detailed description of our system and evaluate its performance. Finally, we discuss the experimental results and suggest areas for improvement.

I. BACKGROUND

Solutions to complex tasks often require the use of multiple techniques for various parts of the problem, and this task is no exception. We were required to implement a localisation algorithm to determine the position of the robot, a probabilistic road map to allow paths to be planned through the space, a navigation algorithm for path planning, a method of exploring the space, and some way of detecting people. We then had to implement a system which would combine all of these separate modules into a single system that would complete the task. In this section we present some background information about techniques that are often used to solve these problems.

A. Localisation

- sensor model, update
- motion model

The aim of localisation is to obtain an estimate of the position of a mobile robot using sensor data [2]. While it is possible to use odometry data to do so, this data is inherently noisy and prone to error. In particular, odometry errors can arise from changes in the surface being traversed and the robot’s weight, among others.

1) *Bayes Filter*: Many advanced localisation techniques are based on the Bayes filter, which uses a belief distribution $bel(x_t)$ to represent the state x_t . The calculation of $bel(x_t)$ at time t is dependent on $bel(x_{t-1})$ at time $t - 1$, the last action u_t , and the last measurement z_t . In the first step, called the prediction step, the prior belief $\bar{bel}(x_t)$ is calculated. This step merges two probability distributions; the prior belief over the previous state x_{t-1} , and the probability of transitioning from that state to the posterior state x_t given that the action u_t was taken. This step does not take into account any measurement taken in the posterior state, predicting based solely on the knowledge of the action taken. In the second step, the measurement update, the posterior belief is calculated by multiplying the prior belief with the probability of being in the posterior state given that the measurement z_t was observed. The result of this multiplication is generally not a probability, and therefore requires normalisation using some constant α .

As there are usually multiple posterior states, x_t is usually a state vector rather than a single state, and so the two steps will be applied multiple times in order to update the belief for each state being considered. The filter is recursive, requiring some idea of the initial belief $bel(x_0)$ at time $t = 0$. The initial belief is either a distribution centred on x_0 , in the case where the initial position is known, or a uniform distribution over the space otherwise. As the Bayes filter is not restricted

Algorithm 1 Bayes filter [3]

```
1: Algorithm Bayes_filter( $bel(x_{t-1}), u_t, z_t$ )
2: for all  $x_t$  do
3:    $\bar{bel}(x_t) = \int P(x_t | u_t, x_{t-1}) bel(x_{t-1}) dx_{t-1}$ 
4:    $bel(x_t) = \alpha P(x_t | z_t) \bar{bel}(x_t)$ 
5: end for
6: return  $bel(x_t)$ 
```

to finite state spaces, it is not possible to implement it for anything other than very simple problems. There are two families of algorithms for localisation, known as *recursive state estimators*, with various properties that permit the use of the Bayes filter in more complex estimation tasks.

2) *Gaussian Filters*: The basic principle of the family of Gaussian filters is the use of multivariate normal distributions, which can be formulated from a mean μ and covariance Σ , to represent belief. As a result, the assumption that the system is a linear Gaussian system is made; the initial belief must be a Gaussian, and both the state transition function and measurement probability must be linear functions. Although Gaussian filters can be extended to non-linear systems, they perform best when the system meets the assumptions made. One of the main advantages of such filters is the computational complexity, which is polynomial with respect to the dimensionality of the state space. The main disadvantage is that Gaussians are unimodal and therefore cannot represent situations in which there are multiple hypotheses; a situation that is often encountered in robotics. Examples include the Kalman filter and the information filter, which are derived from two different ways of representing Gaussians. Both of these filters can be extended to non-linear systems by using a Taylor expansion to produce linear approximations of non-linear functions. Mixtures of Gaussians can also be used to extend the Kalman filter to encompass situations in which multiple hypotheses are required, but each extension increases the complexity of the algorithm. This extensibility is one of the reasons for the popularity of the Kalman filter in state estimation problems. With some extension, the information filter is particularly suited to multi-robot systems where information from multiple sources must be integrated, but issues with complexity have resulted in the Kalman filter becoming the more popular of the two for the majority of problems [3].

3) *Nonparametric Filters*: In contrast to Gaussian filters, nonparametric filters discretise the probability distribution and

do not require assumptions of linearity and Gaussian belief distributions. Instead, the state is approximated by a finite number of values which are taken from the belief state at any given time. The number of these values can be varied, and non-parametric filters converge to the correct state as the number tends towards infinity. Because this family of filters does not impose restrictions on the posterior density, they are useful for problems such as global localisation, which require the state to be represented in a complex form. Global localisation is the problem of determining the position of the robot without knowing its initial position, resulting in global uncertainty and the need for a multimodal belief distribution. Although as a result nonparametric filters are more computationally expensive than Gaussian filters, it is possible to vary the number of values used to represent the belief to suit the problem using *adaptive* techniques. Examples of these types of filters are the histogram and particle filters. The histogram filter decomposes a continuous space into some finite number of regions, each of which is assigned a probability based on the belief. Extensions include using dynamic decomposition techniques to use more coarse representations in regions with lower probability, and the use of *selective updating*; updating only the parts of the space which are deemed important. Particle filters approximate the belief by drawing a number of hypotheses called *particles* randomly from the belief distribution. The most important part of the particle filter is the resampling step, which selects particles from the initial set proportional to an importance factor. This has the effect of concentrating particles in areas of high likelihood, reducing computational power spent in areas which are not relevant. Improving the sampling method and adapting the number of particles based on time and uncertainty lead to more effective and error-resistant particle filters. A property which is particularly useful to us is that particle filters are very easy to implement [3].

4) *The Markov Assumption*: Part of the efficiency of many of these filters comes from their use of the Markov assumption, a property described by the Russian mathematician Andrey Markov. The assumption is that for each of these filters, which is essentially a model of some system, the Markov property holds. While this may not actually be the case, the Bayes filter that forms the base of these filters is robust to the violation of some of the requirements of the property. The Markov property states that the future states of the system do not depend on states in the past. In other words, considering all previous states of the system provide no additional information for predicting future states; all prediction can be done using only the current state. Previous states do not have to be included in calculations, and as a result do not need to be stored, which leads to faster computation and lower storage space requirements.

5) *Sensor & Motion Models*: To actually use any of these filters to solve a localisation problem, it is necessary to model the sensors and motion of the robot in order to calculate probabilities for use in the prediction and measurement updates.

Bayes filter, Kalman filter, particle filter (MCL), small section about mapping—still an active area of research in robotics. Mention SLAM, which has been pretty much solved.

B. Route Planning

PRM (sampling methods, graph search), RRT

Algorithm 2 Basic Monte Carlo Localisation[3]

```

1: Algorithm MCL( $\mathcal{X}_{t-1}, u_t, z_t, m$ )
2:  $\bar{\mathcal{X}}_t = \mathcal{X}_t = \emptyset$ 
3: for  $m = 1$  to  $M$  do
4:    $x_t^{[m]} = \text{sample\_motion\_model}(u_t, x_{t-1}^{[m]})$ 
5:    $w_t^{[m]} = \text{sensor\_model}(z_t, x_t^{[m]}, m)$ 
6:    $\bar{\mathcal{X}}_t = \bar{\mathcal{X}}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$ 
7: end for
8: for  $m = 1$  to  $M$  do
9:   draw  $i$  with probability  $\propto w_t^{[m]}$ 
10:  add  $x_t^{[i]}$  to  $\bar{\mathcal{X}}_t$ 
11: end for
12: return  $\bar{\mathcal{X}}_t$ 

```

Algorithm 3 Probabilistic Road Map Generation

```

1: Algorithm generate_PRM( $map$ )
2:  $V = \text{sample\_vertices}(map)$ 
3: for  $v_c \in V$  do
4:   while  $\text{connections}(v_c) < C$  do
5:      $v_t = \text{get\_closest}(V)$ 
6:     if  $d(v_c, v_t) < D_n$  then
7:       if  $\neg \phi(v_c, v_t) \wedge \gamma(v_c, v_t)$  then
8:         connect( $v_c, v_t$ )
9:       end if
10:    else
11:      end if
12:    end while
13: end for

```

C. Exploration

frontier based techniques

D. Robot Vision

II. DESIGN

A. System Structure

MENTION ALGORITHM COMPLEXITY! brief ROS description, callback based system, finite state automaton

B. Platform

Stuff about the pioneer—available sensors, some data about its size, specifications, our additions to it. Kinect specs. Include

Algorithm 4 Path Flattening

```

1: Algorithm flatten_path( $P, I, map$ )
2: for  $i := 0$  to  $I$  do
3:   for  $j := 0$  to  $|P| - 2$  do
4:      $A \leftarrow \text{newpath}(i)$ 
5:      $B \leftarrow \text{newpath}(i + 1)$ 
6:      $C \leftarrow \text{newpath}(i + 2)$ 
7:     if  $\text{freely\_connected}(map, A, C)$  then
8:        $P \setminus B$ 
9:     end if
10:   end for
11: end for
12: return newpath

```

a picture of the robot with the kinect on it.

III. EXPERIMENTATION

A. PRM

inflated map - show inflated map superimposed onto the original map Redo experiment for sampling methods. short, medium, long path length. Display image of map with one of the routes displayed and show the difference between the sampling methods. Find the optimum route by sampling a massive number of vertices on to the space and then finding a route using that—the flattened path is then the most optimal route, and we compare the other routes to this route for each experiment.

Redid sampling experiments - now have comparison for neighbourhood, threshold and nearestN strategies on 4 different paths, including trying to get into the corridor. 5,10,20,30 maxconn for each strategy, inflation radius of 5 so that you can just about get into the corridor. random vertices:25,50,100,200,400,800, grid step 0.5,1,2,4, cell size 1,2,4,6, target per cell 2,5

REDO THIS ONCE THE ABOVE DATA IS COMPILED Best sampling: (approximate) optimum path found for each path by using grid sampling with a step of 0.2m and 50 max connections and neighbourhood connection strategy. Screenshots available in screenshots dir. Repeated experiments five times for cell and random, once only for grid. 5,10,20,30 maxconnections. Used results from previous experiments on the sampling strategy, but compared the best parameters for each. Also checked whether the corridor which causes issues was accessible when using each strategy as a measure of its ability to populate tight spaces. Not necessarily a good evaluation, since grid may be good on this map by accident, but terrible on others. new_prmlogs contains data.

B. Vision

C. Exploration

Coverage experiments info: cell and grid done for cell size and grid spacing of 10,8,6,4,2,1, with cell repeated three times for each. The fov max distance was 3.5, angle 57, minimum distance 0.3. Started at 2.80,18.97,40 in stage. Max move speed 0.7m/s, max rotation speed 0.4 rad/sec. Ran until the end of the exploration path was reached.

IV. DISCUSSION

A. Performance

B. Potential Improvements

the path generated to do exploration could be improved by using heuristic techniques

C. Conclusions

REFERENCES

- [1] D. Kortenkamp, I. Nourbakhsh, and D. Hinkle, "The 1996 aaai mobile robot competition and exhibition," in *AI Magazine*, vol. 18, 1997.
- [2] W. Burgard, D. Fox, and S. Thrun, "Active mobile robot localization by entropy minimization," in *Proceedings, Second EUROMICRO workshop on Advanced Mobile Robots*, pp. 155–162, 1997.
- [3] S. Thrun, W. Burgard, and D. Fox, *Probabilistic Robotics*. Intelligent Robotics and Autonomous Agents Series, MIT Press, 2005.